

Building a Distributed System with Real-time Constraints

Using concurrent Functional Programming tools

Armando Santos

September 12, 2023 — Copyright © 2023 Well-Typed LLP



Introduction

Networking Team

Armando Santos
Well-Typed

Marcin Szamotulski
Input Output Global

Duncan Coutts
Well-Typed

Neil Davies
PNSol

Peter Thompson
PNSol

What we are responsible for at the
networking team

Cardano Node



Ouroboros algorithm paper	Refining step	Implementation	Testing
Formal Specification	Real time constraints	Architecture & Design	Properties
Threat models	Concurrency	Protocols	Simulation
Non-realistic assumptions	Operation & Performance	Scale	Reliability
Performance objectives		Exceptions & Corner cases	CI

Presentation overview

The Challenge: Real-time Distributed System

Cardano's Vision of a Distributed System

- ▶ It's a distributed system, but
 - ▶ Latency
 - ▶ Scalability
 - ▶ Fault Tolerance
 - ▶ Honest stake assumption
 - ▶ Real-time constraints
 - ▶ Adversarial behavior

Constraints, Assumptions, and Implementation Challenges in Cardano

Faithfulness to Research vs. Real-World Implementation

Ouroboros: A Provably Secure Proof-of-Stake Blockchain Protocol

Aggelos Kiayias^{*} Alexander Russell[†] Bernardo David[‡] Roman Oliynykov[§]

July 20, 2019

Abstract

We present “Ouroboros,” the first blockchain protocol based on *proof of stake* with rigorous security guarantees. We establish security properties for the protocol comparable to those achieved by the Bitcoin blockchain protocol. As the protocol provides a “proof of stake” blockchain discipline, it offers qualitative efficiency advantages over blockchains based on *proof of physical resources* (e.g., proof of work). We also present a novel reward mechanism for incentivizing proof of state protocols and we prove that, given this mechanism, honest behavior is an approximate Nash equilibrium, thus neutralizing attacks such as selfish mining. We also present initial evidence of the practicality of our protocol in real world settings by providing experimental results on transaction confirmation and processing.

Figure: Ouroboros paper

```
(- HLINT ignore "Fdisk concatenates" -)
(- HLINT ignore "Redundant <code>" -)
(- HLINT ignore "Use fewer imports" -)

runMainIO
  $ PartialNodeConfiguration
  = IO (|
    PartialConfig = do
      installSigTermHandler
      Crypto.cryptoInit
      configWithPc = parseNodeConfigurationPc = getLast $ pcConfigFile cmdPc
      nc = case makeNodeConfiguration $ defaultPartialNodeConfigurations $ configWithPc $ cmdPc of
        Left err => error $ "Error in creating the NodeConfiguration: " ++ err
        Right nc' => return nc'
      putStrLn "Node configuration: " ++ show nc
      case shelleyWRFFile $ ncProtocolFiles nc of
        Just wrfp = do vrf = runExceptT $ checkWRFFilePermissions (File wrfp)
                      case vrf of
                        Left err => putStrLn ("Error in checking WRF file permissions: " ++ err)
                        Right () => pure ()
        Nothing = pure ()
      eitherSomeProtocol = runExceptT $ mkConsensusProtocol
        (mkProtocolConfig nc)
        -- TSO: Convert ncProtocolFiles to Maybe as relay nodes
        -- don't need these.
        (Just $ ncProtocolFiles nc)

      p @ SomeConsensusProtocol =
      case mkSomeProtocol of
        Left err => putStrLn ("Error in creating consensus protocol: " ++ err)
        Right p' = pure p'

      let networkMagic :: Api.NetworkMagic =
        case p of
          SomeConsensusProtocol _ runP =
            let ProtocolsInfo { plateauConfig } = fst $ Api.protocolInfo @IO runP
            in getNetworkMagic $ Consensus.configBlock plateauConfig

      case d @ SomeConsensusProtocol blockType runP =
        handleHeaderWithTracers cmdPc nc p networkMagic blockType runP

      -- Used here to ensure that the main thread throws an async exception on
      -- receiving a SIGTERM signal.
      installSigTermHandler = IO (|
        installSigTermHandler = do
          installHandler
          -- Similar implementation to the RTS's handling of SIGINT (see GHC's
          -- runtime/glasgow-haskell-prim/ghc/Blab/master/Libraries/base/GHC/TopHandler.hs),
          -- runMainIO ::##################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################################### ######## ####### Input Output logo


Well-Typed


```

Need to fill the gap

- ▶ Translate formal specification into code
- ▶ Highly concurrent logic
- ▶ High assurance needs (i.e. as verifiable and robust as possible)
- ▶ Clear code structures, predictable behavior

Strongly Statically Typed Purely Functional Programming

Haskell



Figure: Source

But:

- ▶ Memory consumption
- ▶ Debugging can be tricky

In a nutshell



Figure: Source

Tools & Techniques

Software Transactional Memory (STM)

- ▶ Compositional and modular concurrency
- ▶ Isolated from IO side effects

Abstraction & Libraries

- ▶ First to finish vs Last to finish STM semantics
- ▶ Typed-protocols
 - ▶ Type safe
 - ▶ Session types
 - ▶ Deadlock free!

```
... LastToFinish synchronization. It is the multiplicative semigroup of
... the linear semiring (https://www.akkademy.com/linear-semiring) for which addition is
... given by 'FirstToFinish'.
...
... This is similar to 'Ap' (use 'LastToFinishM') in the sense that it will wait
... for all monadic effects, but unlike 'Ap' it will not combine all results
... using a monoid instance, but rather it will return the last one. Also unlike
... 'Ap' it does not have a monoid unit.
...
... | Read all the TMVar's and return the one that was filled last.
...
... | Read all the TMVar's and return the one that was filled last.
...
... readAllTMVars :: Monad m => Nonempty (TMVar m) --> STM a
... readAllTMVars = runlastToFinish
...   foldM (lastToFinish, readTMVar)
...   ...
...   ;`from "semigroupoids" package or use 'foldM' and 'Fmap'
...   ;`from "base"
...
newtype LastToFinish m a = LastToFinish { runLastToFinish :: m a }
deriving newtype Functor
deriving
  Generic
  Foldable
  Traversable
  Alternative
  Monoid
  MonoidPlus
  Traversable
deriving Foldable via (Ap m)
instance MonoidPlus m => Semigroup (LastToFinish m a) where
  (m1 :> m2) = LastToFinish m a -> LastToFinish m a -> LastToFinish m a
  LastToFinish left #> LastToFinish right = LastToFinish $ do
    m1 <- left
    m2 <- right
    case m of
      Left x = right
      Right y = left
      m     = m2
```

3.13.2 State machine

Agency	
Client has Agency	SIdle
Server has Agency	SBusy

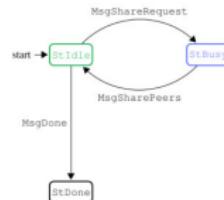


Figure 3.10: State machine of the peer sharing protocol.

Testing approach

- ▶ Property based testing vs unit tests
 - ▶ Input random generation;
 - ▶ Shrinking;
 - ▶ Reproducibility;
 - ▶ Coverage checks.
- ▶ Simulation
 - ▶ Early detection of critical races;
 - ▶ Simulation of rare edge cases;
 - ▶ Mocking and error injection;
 - ▶ Simulate time passing;
 - ▶ Looking for different schedules.

Highlights from the Networking Team

Highlights

Good:

- ▶ Found several tricky edge case bugs
- ▶ Network is running for years non-stop
- ▶ Simulate years of total time per week!

But:

- ▶ We are far from perfect, we still fix and find small bugs every month related to networking

We try our best:

Our CI test suite runs on average between 1 and 5 hours of simulated time per test per input per PR per OS. So:

- ▶ Assuming around 100 tests in our test suite and
- ▶ each test generates 100 random inputs;
- ▶ Assuming 3 PRs per week and testing on Windows, OSX and Linux;
- ▶ Results on around **225 000 hours** of simulated time per week!

Takeaways

Takeaways

- ▶ Complex systems spans performance characteristics we can not control;
- ▶ Functional Programming, namely Haskell and its concurrency tools helps us manage complexity;
- ▶ We do our best in searching through all state space efficiently;
- ▶ Well-founded confidence;
- ▶ Progress is achievable.

Thank you!